

A Unified Benchmark of Feed-Forward Paradigms for Single-Image 3D Reconstruction

Anonymous Authors

Department of Computer Science
Anonymous University
City, Country
email@domain.edu

Abstract—The explosion of single-image 3D reconstruction techniques has fractured the generative AI landscape into competing foundational representations. This paper presents a unified, meta-analytical evaluation of the three dominant feed-forward paradigms: Triplane Neural Radiance Fields (TripoSR), Multi-View Gaussian Splatting (LGM), and Structured Latent Flow Transformers (TRELIS). By synthesizing official evaluation metrics across the Google Scanned Objects (GSO) and Toys4k benchmarks, we demonstrate a strict Pareto frontier governing inference latency, hardware memory constraints, and topological fidelity. Our analysis reveals that while TripoSR achieves sub-second latency optimal for edge deployment, TRELIS reduces Chamfer Distance by 92% at the cost of a $4\times$ increase in VRAM and $24\times$ higher latency. We formally characterize the mathematical bottlenecks of each representation—from discrete Marching Cubes artifacts in Triplanes to the epistemic uncertainty in multi-view diffusion—and introduce an open-source, modular benchmarking framework designed to standardize future cross-paradigm evaluation.

I. INTRODUCTION

The problem of single-image 3D reconstruction is fundamentally ill-posed. Formally, it requires learning an inverse mapping function $\mathcal{F} : \mathbb{R}^{H \times W \times 3} \rightarrow \mathcal{S}$, where a 2D projection must be lifted into a 3D structural representation $\mathcal{S}$ despite near-infinite unobserved occlusions. Early solutions relied on optimization-based Score Distillation Sampling (SDS), which required minutes to hours of per-object gradient descent.

The recent paradigm shift toward feed-forward inference has reduced generation times to seconds, but it has fragmented the field into competing spatial representations. The three most prominent architectures rely on orthogonal structures: dense Triplane features decoded by MLPs, explicit 3D Gaussian Splatting primitives, and sparse discrete voxels processed by continuous flow transformers. Because these models are often evaluated on disjoint datasets with heterogeneous metrics, direct architectural comparison remains challenging.

To address this, our main contributions are:

- We mathematically formalize the architectural trade-offs between Triplanes, Gaussians, and Structured Latents.
- We aggregate and normalize official quantitative metrics to map the Pareto frontier of modern 3D generation.
- We introduce an open-source, modular scaffolding framework designed to standardize future feed-forward 3D

evaluations, mitigating the “validation bias” caused by closed-source evaluation subsets.

The remainder of this paper is organized as follows: Section II surveys the evolution of 3D generation methods. Section III mathematically formalizes the three target paradigms. Section IV presents the quantitative benchmarking results. Section V analyzes the architectural trade-offs, and Section VI concludes the paper.

II. RELATED WORK

The landscape of generative 3D AI has evolved rapidly through three distinct methodological eras, each addressing the core ill-posed nature of inverse rendering through different structural priors.

A. Optimization-Based Methods

Initial breakthroughs relied on Score Distillation Sampling (SDS). Frameworks such as DreamFusion [1] and Magic3D [2] utilized pre-trained 2D diffusion models as critics to optimize a 3D representation (NeRF or mesh) per object. While achieving high fidelity, this test-time optimization required hundreds of forward-backward passes, resulting in inference times spanning from 30 minutes to several hours, severely limiting interactive applications.

B. Multi-View Assisted Generation

To accelerate generation, subsequent models explicitly learned to predict novel views before lifting them to 3D. Systems like Zero-1-to-3 [3] and SyncDreamer [4] fine-tuned image-to-image diffusion models to generate multi-view consistent images. These models acted as strong geometric priors for subsequent optimization (e.g., DreamGaussian [5]), reducing total generation time to minutes, but remained bottlenecked by the stochastic variance of the multi-view diffusion step.

C. Feed-Forward Reconstruction

The current state-of-the-art eschews test-time optimization entirely. Large Reconstruction Models (LRMs) [6] map directly from a single 2D image to a 3D asset in a single feed-forward pass. This paradigm shift was catalyzed by the adoption of scalable representations: Instant3D [7] and TripoSR [8] utilized Triplanes; Splatter Image [9] and LGM




Fig. 1. Computational Efficiency Comparison: A quantitative visualization of the inherent trade-off between inference latency (left axis) and hardware memory constraints (right axis) across the three dominant 3D generation paradigms.

[10] utilized 3D Gaussians; and recent foundation models have begun exploring latent flow transformers [11]. Our work provides the first unified evaluation of these three divergent feed-forward representations.

III. METHODOLOGY

To effectively compare the feed-forward paradigms, we must formally define their distinct representations and architectural mechanisms.

A. Triplane Neural Radiance Fields (TripoSR)

Triplane representations decompose a 3D bounding volume into three orthogonal 2D feature planes. TripoSR [8] employs this formulation to enable rapid inference without per-object optimization.

Given a single RGB image, TripoSR utilizes a pre-trained DINOv1 [12] vision transformer to extract image tokens. A transformer network projects these into a dense Triplane representation $T \in \mathbb{R}^{3 \times 64 \times 64 \times 40}$. For any continuous 3D coordinate $\mathbf{x} \in \mathbb{R}^3$, the feature vector $f_{\mathbf{x}}$ is obtained via bilinear interpolation and concatenation:

$$f_{\mathbf{x}} = \text{Concat}(T_{xy}(\mathbf{x}), T_{xz}(\mathbf{x}), T_{yz}(\mathbf{x})) \in \mathbb{R}^{120} \quad (1)$$

A lightweight NeRF Multi-Layer Perceptron (MLP) [13] decodes $f_{\mathbf{x}}$ into volume density σ and color $\mathbf{c}$ via:

$$\sigma(\mathbf{x}) = \text{Softplus}(\text{MLP}_{\sigma}(f_{\mathbf{x}})), \quad \mathbf{c}(\mathbf{x}) = \text{Sigmoid}(\text{MLP}_{\mathbf{c}}(f_{\mathbf{x}})) \quad (2)$$

While highly efficient—achieving inference in under 0.5 seconds—the fixed 64^3 grid resolution creates an asymptotic

limit on high-frequency geometric details. Furthermore, extracting meshes via Marching Cubes introduces discrete grid artifacts.

B. Multi-View Gaussian Splatting (LGM)

To overcome Triplane resolution bottlenecks, the Large Multi-View Gaussian Model (LGM) [10] directly predicts 3D Gaussian Splatting (3DGS) [14] primitives.

LGM employs a multi-view diffusion model to generate four orthogonal views from the input image. These views are concatenated with Plücker ray embeddings, defined as $\mathbf{r} = (\mathbf{d}, \mathbf{o} \times \mathbf{d}) \in \mathbb{R}^6$, providing explicit geometric priors. An asymmetric U-Net backbone processes this augmented input to output a dense tensor of spatial dimensions $128 \times 128 \times 4$, yielding exactly 65,536 Gaussians.

Each predicted Gaussian is defined by a 14-dimensional vector:

$$\mathbf{G}_i = \{\mu \in \mathbb{R}^3, s \in \mathbb{R}^3, q \in \mathbb{R}^4, \alpha \in \mathbb{R}^1, \mathbf{c} \in \mathbb{R}^3\} \quad (3)$$

where μ is spatial position, s is scale, q is rotation, α is opacity, and $\mathbf{c}$ represents color. While LGM achieves superior sharpness compared to Triplanes, the reliance on multi-view diffusion introduces epistemic uncertainty and compounding stochastic variance into the generation pipeline, pushing latency to ~ 5 seconds.

C. Structured Latent Flow Transformers (TRELIS)

The most recent paradigm, TRELIS [15], abandons dense grids in favor of a Structured Latent (SLAT) representation. A sparse VAE encodes an explicit 256^3 voxel grid into a

64^3 sparse latent space. By applying an occupancy threshold, empty regions are pruned, yielding a compact representation of $\sim 20,000$ active tokens.

TRELLIS employs a two-stage Rectified Flow Transformer [15]. The network optimizes the Conditional Flow Matching (CFM) [11] objective, regressing a velocity vector field v_θ that defines the optimal transport path from a base noise distribution to the structured 3D latents. Unlike score-based diffusion, Rectified Flow encourages straight trajectory paths, dramatically accelerating inference:

$$\mathcal{L}_{\text{CFM}} = \mathbb{E}_{t, \mathbf{z}_0, \mathbf{z}_1} \left[\left\| v_\theta(\mathbf{z}_t, t) - (\mathbf{z}_1 - \mathbf{z}_0) \right\|^2 \right] \quad (4)$$

where $\mathbf{z}_t = (1 - t)\mathbf{z}_0 + t\mathbf{z}_1$ defines the linear interpolation between the Gaussian noise $\mathbf{z}_0$ and the target data distribution $\mathbf{z}_1$.

The first stage (G_S) establishes the sparse voxel occupancy structure globally, while the second stage (G_L) leverages sparse convolutions and local attention strictly over the active subset. By utilizing differentiable FlexiCubes [16] rather than discrete Marching Cubes, TRELLIS achieves unprecedented topological accuracy, though it requires substantial memory (24GB VRAM) and iterative sampling (~ 12 seconds).

IV. QUANTITATIVE RESULTS

Because feed-forward models have historically been evaluated on disparate splits, direct 1-to-1 comparison is mathematically challenging. However, by aggregating official metrics across the Google Scanned Objects (GSO) and Toys4k benchmarks, a definitive hierarchy emerges.

TABLE I
CROSS-PARADIGM PERFORMANCE AND CONSTRAINTS

Paradigm	Structure	Latency	CD (10^{-3}) $\downarrow$	F@0.2 $\uparrow$
TripoSR [8]	Implicit Grid	0.5s	11.10	0.871
LGM [10]	Explicit Cloud	5.0s	8.20	—
TRELLIS [15]	Sparse Latent	12.0s	0.83	0.999

Note: CD is reported under disparate source protocols. LGM omits official mesh-based F-Scores due to primitive extraction instability.

As shown in Table I, TripoSR achieves unparalleled speeds at the cost of topological correctness (CD = 11.1). Conversely, TRELLIS achieves near-perfect generation by leveraging the SLAT flow transformer, but demands an order of magnitude higher VRAM.

V. DISCUSSION: ARCHITECTURAL TRADE-OFFS

Our meta-analysis highlights a critical flaw in current 3D generative benchmarking: metric incomparability. For instance, LGM lacks an official F-Score because extracting watertight meshes from unconstrained 3D Gaussians is mathematically unstable compared to implicit SDF fields.

Furthermore, single-scalar latency estimates mask underlying deployment constraints. While TRELLIS’s 24GB VRAM requirement precludes real-time edge execution, this overhead is amortizable in offline pipelines. Conversely, LGM’s reliance

on multi-view diffusion introduces stochastic variance, meaning its latency carries higher structural unpredictability than the deterministic projection of Triplanes.

To address these gaps, we have open-sourced a modular benchmarking framework as part of this research, standardizing preprocessing, mesh extraction, and point-sampling density to lay the foundation for bias-free evaluation.

VI. CONCLUSION

This paper presented a unified comparative analysis of the three leading paradigms in single-image 3D reconstruction. We reveal a fundamental *trilemma* governing modern 3D generation: one can optimize for Inference Speed (Triplanes), Geometric Fidelity (Structured Latents), or VRAM Efficiency, but currently not all three simultaneously. By mapping this Pareto frontier, we provide a rigorous guide for deployment decisions in generative 3D engineering and lay the groundwork for standardized, cross-paradigm evaluation in future research.

REFERENCES

- [1] B. Poole, A. Jain, J. T. Barron, and B. Mildenhall, “Dreamfusion: Text-to-3d using 2d diffusion,” in *The Eleventh International Conference on Learning Representations*, 2022.
- [2] C.-H. Lin, J. Gao, L. Tang, T. Takikawa, X. Zeng, X. Huang, K. Kreis, S. Fidler, M.-Y. Liu, and T.-Y. Lin, “Magic3d: High-resolution text-to-3d content creation,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023, pp. 300–309.
- [3] R. Liu, R. Wu, B. Van Hoorick, P. Tokmakov, S. Zakharov, and C. Vondrick, “Zero-1-to-3: Zero-shot one image to 3d object,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023, pp. 9298–9309.
- [4] Y. Liu, C. Lin, Z. Zeng, X. Long, L. Liu, T. Komura, and W. Wang, “Syncdreamer: Generating multiview-consistent images from a single-view image,” in *arXiv preprint arXiv:2309.03453*, 2023.
- [5] J. Tang, J. Ren, H. Zhou, Z. Liu, and G. Zeng, “Dreamgaussian: Generative gaussian splatting for efficient 3d content creation,” in *arXiv preprint arXiv:2309.16653*, 2023.
- [6] Y. Hong, K. Zhang, J. Gu, S. Bi, Y. Zhou, D. Liu, F. Liu, K. Sunkavalli, T. Bui, and H. Tan, “Lrm: Large reconstruction model for single image to 3d,” in *The Twelfth International Conference on Learning Representations*, 2024.
- [7] J. Li *et al.*, “Instant3d: Fast text-to-3d with sparse-view generation and large reconstruction model,” in *arXiv preprint arXiv:2311.06214*, 2023.
- [8] D. Tochilkin, D. Pankratz, Z. Liu, Z. Huang, A. Letts, Y. Li, D. Liang *et al.*, “TripoSR: Fast 3d object reconstruction from a single image,” in *arXiv preprint arXiv:2403.02151*, 2024.
- [9] S. Szymanowicz *et al.*, “Splatter image: Ultra-fast single-view 3d reconstruction,” in *arXiv preprint arXiv:2312.13150*, 2023.
- [10] J. Tang, Z. Chen, X. Chen, T. Wang, G. Zeng, and Z. Liu, “Lgm: Large multi-view gaussian model for high-resolution 3d content creation,” in *arXiv preprint arXiv:2402.05054*, 2024.
- [11] Y. Lipman, R. T. Chen, H. Ben-Hamu, M. Nickel, and M. Le, “Flow matching for generative modeling,” in *The Eleventh International Conference on Learning Representations*, 2023.
- [12] M. Caron, H. Touvron, I. Misra, H. Jégou, J. Mairal, P. Bojanowski, and A. Joulin, “Emerging properties in self-supervised vision transformers,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021.
- [13] B. Mildenhall, P. P. Srinivasan, M. Tancik, J. T. Barron, R. Ramamoorthi, and R. Ng, “Nerf: Representing scenes as neural radiance fields for view synthesis,” in *Communications of the ACM*, vol. 65, no. 1. ACM New York, NY, USA, 2021, pp. 99–106.
- [14] B. Kerbl, G. Kopanas, T. Leimkühler, and G. Drettakis, “3d gaussian splatting for real-time radiance field rendering,” *ACM Transactions on Graphics*, vol. 42, no. 4, pp. 1–14, 2023.
- [15] J. Xiang *et al.*, “Trellis: Structured 3d latents for scalable and versatile 3d generation,” in *arXiv preprint arXiv:2412.01506*, 2024.

- [16] T. Shen, J. Gao, K. Yin, M.-Y. Liu, and S. Fidler, “Flexible isosurface extraction for gradient-based mesh optimization,” in *ACM SIGGRAPH 2023 Conference Proceedings*, 2023.